

Capstone Project

February 7, 2021

1 Capstone Project

1.1 Image classifier for the SVHN dataset

1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
import random
from scipy.io import loadmat
from matplotlib.pyplot import figure, imshow, axis,subplots
import numpy as np
```



For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]: *# Run this cell to load the dataset*

```
train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both train and test are dictionaries with keys X and y for the input images and labels respectively.

1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

In [37]: *# extract data and labels*

```
x_train = train['X']
y_train = train['y'] - 1
x_test = test['X']
y_test = test['y'] - 1      #convert to 0-indexed
```

```

In [39]: # display 5 random figures from training set and 5 random figures from test set
train_sample_size = 5
test_sample_size = 5
train_size = x_train.shape[3]
test_size = x_test.shape[3]
random_train_index = random.sample(range(0, train_size), train_sample_size)
random_test_index = random.sample(range(0, test_size), test_sample_size)
_, train_axis = subplots(1, train_sample_size)
_, test_axis = subplots(1, test_sample_size)
# display
for i in range(0, train_sample_size):
    train_axis[i].imshow(x_train[:, :, :, random_train_index[i]])
    train_axis[i].axis('off')
    train_axis[i].set_title(str(y_train[random_train_index[i]]+1))
for i in range(0, test_sample_size):
    test_axis[i].imshow(x_test[:, :, :, random_test_index[i]])
    test_axis[i].axis('off')
    test_axis[i].set_title(str(y_test[random_test_index[i]]+1))

```



```

In [43]: # compute average along color channel.
x_train = np.mean(x_train, axis=2, keepdims=True)
x_test = np.mean(x_test, axis=2, keepdims=True)

```

```

In [44]: # display 5 random figures from training set and 5 random figures from test set
train_sample_size = 5

```

```

test_sample_size = 5
train_size = x_train.shape[3]
test_size = x_test.shape[3]
random_train_index = random.sample(range(0, train_size), train_sample_size)
random_test_index = random.sample(range(0, test_size), test_sample_size)
_, train_axis = subplots(1, train_sample_size)
_, test_axis = subplots(1, test_sample_size)
train_axis[0].set_axis_off()
# display
for i in range(0, train_sample_size):
    train_axis[i].imshow(x_train[:, :, 0, random_train_index[i]], cmap="gray")
    train_axis[i].axis('off')
    train_axis[i].set_title(str(y_train[random_train_index[i]]))
for i in range(0, test_sample_size):
    test_axis[i].imshow(x_test[:, :, 0, random_test_index[i]], cmap="gray")
    test_axis[i].axis('off')
    test_axis[i].set_title(str(y_test[random_test_index[i]]))

```



```

In [45]: # move the input data axis to the first place
x_train = np.moveaxis(x_train, 3, 0)
x_test = np.moveaxis(x_test, 3, 0)

```

1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.

- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the `summary()` method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a `ModelCheckpoint` callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [8]: # tf imports
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Flatten
        from tensorflow.keras.optimizers import Adam
        from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
In [9]: def build_model(input_shape):
        model = Sequential([
            Flatten(input_shape=input_shape, name='flatten_1'),
            Dense(512, activation='relu', name='dense_1'),
            Dense(256, activation='relu', name='dense_2'),
            Dense(10, activation='softmax', name='output')
        ])
        #compile
        model.compile(optimizer='adam', loss=tf.keras.losses.SparseCategoricalCrossentropy)
        return model
```

```
In [10]: #get input shape from one image
         input_shape = x_train.shape[1:]
         model = build_model(input_shape)
         model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dense_2 (Dense)	(None, 256)	131328
output (Dense)	(None, 10)	2570

=====
Total params: 658,698

Trainable params: 658,698
Non-trainable params: 0

In [11]: *# save callbacks*

```
best_ckpt_path = 'model_checkpoint_best/checkpoint'  
best_ckpt = ModelCheckpoint(filepath= best_ckpt_path,  
                             save_weights_only=True,  
                             monitor='val_accuracy',  
                             save_best_only=True,  
                             verbose=0)
```

stopping callbacks

```
early_stop = EarlyStopping(patience=2)
```

In [13]: *# training phase*

```
history= model.fit(x_train,y_train, epochs=30,validation_split=0.15, batch_size=64,verbose=1,  
                  callbacks=[best_ckpt,early_stop])
```

Train on 62268 samples, validate on 10989 samples

Epoch 1/30

62268/62268 - 33s - loss: 4.8971 - accuracy: 0.1873 - val_loss: 2.0356 - val_accuracy: 0.3471

Epoch 2/30

62268/62268 - 32s - loss: 1.7459 - accuracy: 0.4438 - val_loss: 1.6254 - val_accuracy: 0.4919

Epoch 3/30

62268/62268 - 32s - loss: 1.4212 - accuracy: 0.5527 - val_loss: 1.2697 - val_accuracy: 0.5978

Epoch 4/30

62268/62268 - 32s - loss: 1.3007 - accuracy: 0.5949 - val_loss: 1.1984 - val_accuracy: 0.6308

Epoch 5/30

62268/62268 - 32s - loss: 1.2537 - accuracy: 0.6135 - val_loss: 1.2533 - val_accuracy: 0.6141

Epoch 6/30

62268/62268 - 32s - loss: 1.2381 - accuracy: 0.6158 - val_loss: 1.2935 - val_accuracy: 0.5973

In [14]: *import matplotlib.pyplot as plt*

```
fig = plt.figure(figsize=(12, 5))
```

```
fig.add_subplot(121)
```

```
plt.plot(history.history['loss'])
```

```
plt.plot(history.history['val_loss'])
```

```
plt.title('loss vs. epochs')
```

```
plt.ylabel('Loss')
```

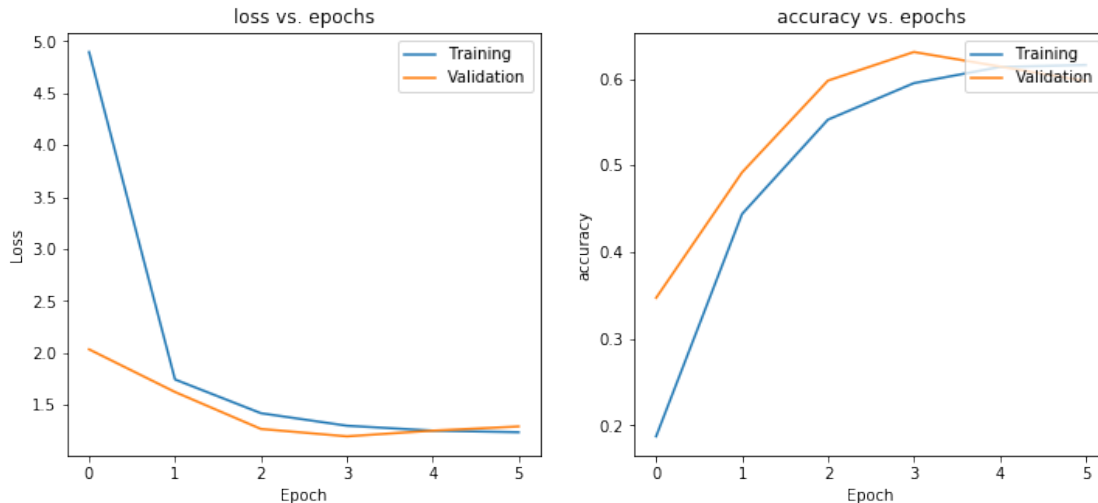
```
plt.xlabel('Epoch')
```

```
plt.legend(['Training', 'Validation'], loc='upper right')
```

```
fig.add_subplot(122)
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy vs. epochs')
plt.ylabel('accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```



```
In [28]: # evaluation on test set
         loss, accuracy = model.evaluate(x_test, y_test, verbose=2)
         print("MLP loss on test set is {}".format(loss))
         print("MLP accuracy on test set is {}".format(accuracy))
```

```
26032/1 - 6s - loss: 1.1501 - accuracy: 0.5900
MLP loss on test set is 1.4010111775292746
MLP accuracy on test set is 0.5899661779403687
```

1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.

- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [19]: from tensorflow.keras.layers import Conv2D, BatchNormalization, MaxPooling2D
def build_CNN_model(input_shape):
    model = Sequential([
        Conv2D(filters=4, kernel_size=(3,3),activation='relu', padding='same', kernel.
        BatchNormalization(), #channel last
        MaxPooling2D(pool_size=(2,2)),
        Conv2D(filters=8, kernel_size=(3,3),activation='relu', padding='same', kernel.
        BatchNormalization(), #channel last
        MaxPooling2D(pool_size=(2,2)),
        Flatten(),
        Dense(100, activation='relu',name='dense_1'),
        Dense(10, activation='softmax',name='output')
    ])
    model.compile(optimizer='adam', loss = tf.keras.losses.SparseCategoricalCrossentr
    return model
```

```
In [20]: #get input shape from one image
input_shape = x_train.shape[1:]
cnn_model = build_CNN_model(input_shape)
cnn_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 4)	40
batch_normalization_4 (Batch Normalization)	(None, 32, 32, 4)	16
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 4)	0
conv2d_2 (Conv2D)	(None, 16, 16, 8)	296
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 8)	32
max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 8)	0
flatten_2 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 100)	51300


```

output (Dense)                (None, 10)                1010
=====
Total params: 52,694
Trainable params: 52,670
Non-trainable params: 24
-----

```

```
In [21]: # save callbacks
```

```

cnn_best_ckpt_path = 'cnn_model_checkpoint_best/checkpoint'
cnn_best_ckpt = ModelCheckpoint(filepath= cnn_best_ckpt_path,
                                save_weights_only=True,
                                monitor='val_accuracy',
                                save_best_only=True,
                                verbose=0)

```

```
# stopping callbacks
```

```
early_stop = EarlyStopping(patience=2)
```

```
In [22]: # training phase
```

```

cnn_history= cnn_model.fit(x_train,y_train, epochs=30, validation_split=0.15, batch_s
                        callbacks=[cnn_best_ckpt,early_stop])

```

Train on 62268 samples, validate on 10989 samples

Epoch 1/30

62268/62268 - 178s - loss: 1.0272 - accuracy: 0.6696 - val_loss: 0.6577 - val_accuracy: 0.8006

Epoch 2/30

62268/62268 - 181s - loss: 0.5717 - accuracy: 0.8253 - val_loss: 0.5674 - val_accuracy: 0.8306

Epoch 3/30

62268/62268 - 178s - loss: 0.4819 - accuracy: 0.8524 - val_loss: 0.5369 - val_accuracy: 0.8418

Epoch 4/30

62268/62268 - 169s - loss: 0.4226 - accuracy: 0.8694 - val_loss: 0.5283 - val_accuracy: 0.8435

Epoch 5/30

62268/62268 - 168s - loss: 0.3793 - accuracy: 0.8834 - val_loss: 0.5201 - val_accuracy: 0.8488

Epoch 6/30

62268/62268 - 169s - loss: 0.3413 - accuracy: 0.8952 - val_loss: 0.5284 - val_accuracy: 0.8486

Epoch 7/30

62268/62268 - 168s - loss: 0.3133 - accuracy: 0.9043 - val_loss: 0.5239 - val_accuracy: 0.8537

```
In [23]: fig = plt.figure(figsize=(12, 5))
```

```
fig.add_subplot(121)
```

```

plt.plot(cnn_history.history['loss'])
plt.plot(cnn_history.history['val_loss'])
plt.title('CNN loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')

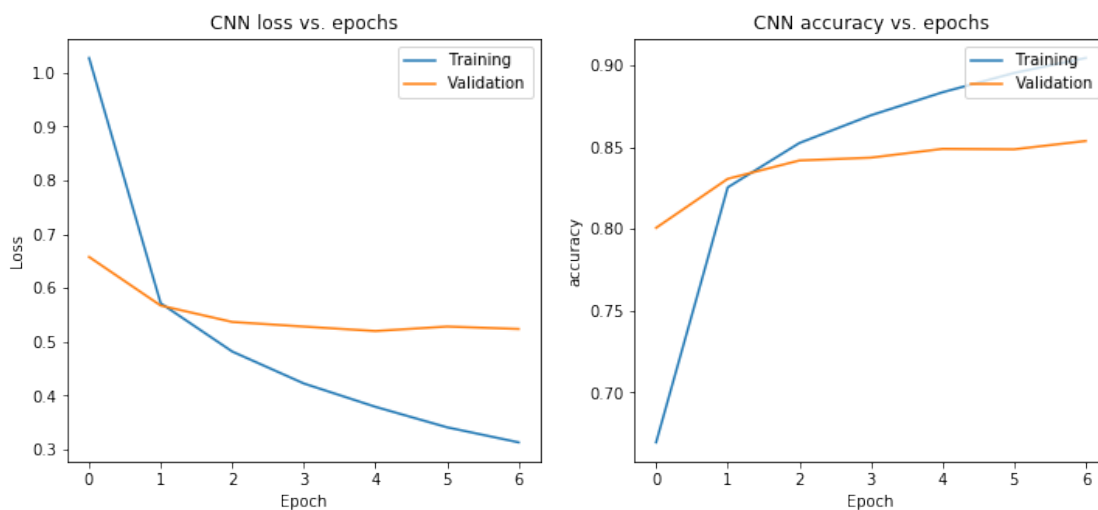
```

```
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)

plt.plot(cnn_history.history['accuracy'])
plt.plot(cnn_history.history['val_accuracy'])
plt.title('CNN accuracy vs. epochs')
plt.ylabel('accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```



```
In [29]: # evaluation on test set
         loss, accuracy = cnn_model.evaluate(x_test, y_test, verbose=2)
         print("CNN loss on test set is {}".format(loss))
         print("CNN accuracy on test set is {}".format(accuracy))
```

26032/1 - 25s - loss: 0.5315 - accuracy: 0.8374

CNN loss on test set is 0.6039436107436132

CNN accuracy on test set is 0.8373540043830872

1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
In [31]: #load best MLP
load_MLP_model = build_model(input_shape)
load_MLP_model.load_weights(best_ckpt_path)
load_MLP_model.evaluate(x_test,y_test, verbose=2)
```

26032/1 - 6s - loss: 1.0726 - accuracy: 0.5943

```
Out[31]: [1.376613517370265, 0.594307]
```


```
In [33]: #load best CNN
load_CNN_model = build_CNN_model(input_shape)
load_CNN_model.load_weights(cnn_best_ckpt_path)
load_CNN_model.evaluate(x_test,y_test, verbose=2)
```

26032/1 - 23s - loss: 0.5315 - accuracy: 0.8374

```
Out[33]: [0.6039436107436132, 0.837354]
```

```
In [54]: #random selection
sample_size =5
random_test_index_pred = random.sample(range(0, test_size), sample_size)
_, pred_axis = subplots(1,sample_size)

for i in range(0, sample_size):
    pred_axis[i].imshow(x_test[random_test_index_pred[i],:,:,0],cmap="gray")
    pred_axis[i].axis('off')
    pred_axis[i].set_title(str(y_test[random_test_index_pred[i]]+1))
```



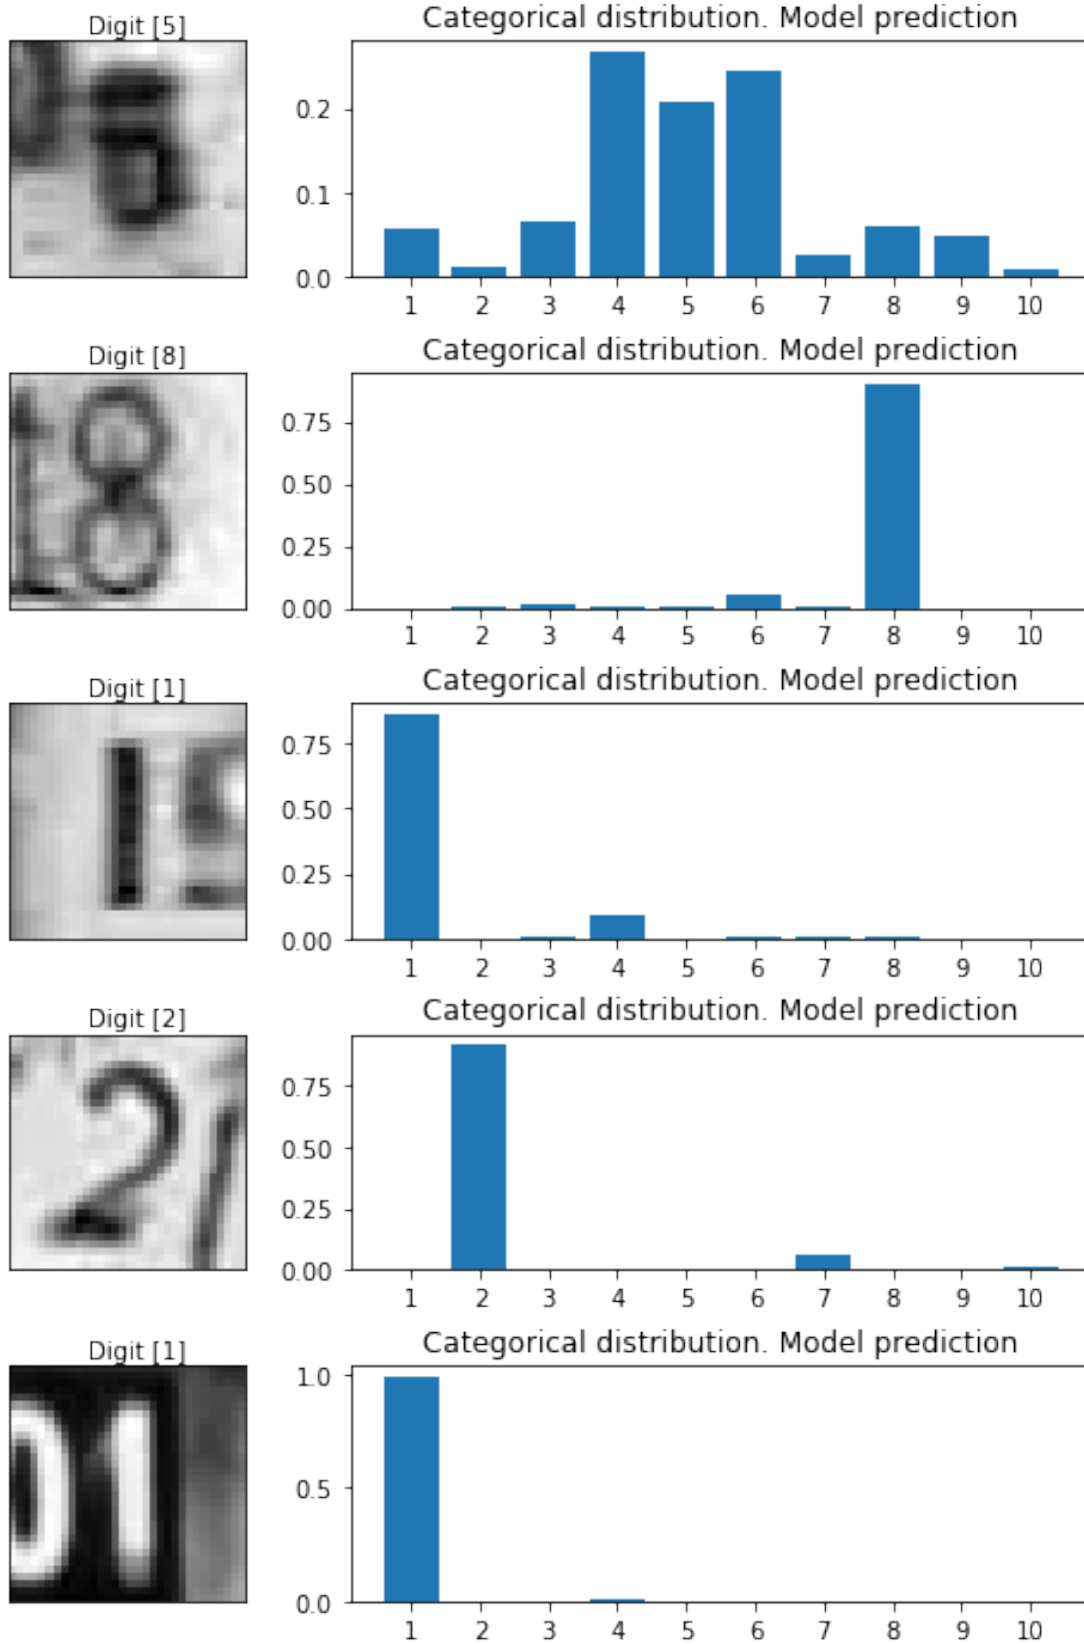
```
In [79]: #display MLP rediction result
pred_image_set = x_test[random_test_index_pred,:,:,0]
pred_label_set = y_test[random_test_index_pred,:]
result = load_MLP_model.predict(pred_image_set)
predictions = np.argmax(result,axis=1)+1
fig, axes = plt.subplots(5, 2, figsize=(10, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)
```

```

for i, (prediction, image, label) in enumerate(zip(result, pred_image_set, pred_label)):
    axes[i, 0].imshow(np.squeeze(image), cmap="gray")
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Digit {label+1}')
    axes[i, 1].bar(np.arange(1,11), prediction)
    axes[i, 1].set_xticks(np.arange(1,11))
    axes[i, 1].set_title("Categorical distribution. Model prediction")

plt.show()
print("predictions:{}".format(predictions))

```




```
predictions:[4 8 1 2 1]
```

```
In [80]: #display CNN rediction result
```

```
cnn_result = load_CNN_model.predict(pred_image_set)
cnn_prediction = np.argmax(cnn_result,axis=1)+1
fig, axes = plt.subplots(5, 2, figsize=(10, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)
```

```
for i, (prediction, image, label) in enumerate(zip(cnn_result, pred_image_set, pred_label)):
    axes[i, 0].imshow(np.squeeze(image),cmap="gray")
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Digit {label+1}')
    axes[i, 1].bar(np.arange(1,11), prediction)
    axes[i, 1].set_xticks(np.arange(1,11))
    axes[i, 1].set_title("Categorical distribution. Model prediction")
```

```
plt.show()
print("predictions:{}".format(cnn_prediction))
```



```
predictions:[5 8 1 2 1]
```